

Bil 470 / YAP 470

Introduction to Machine Learning (Yapay Öğrenme)

Batuhan Bardak

Lecture 4: Evaluation metrics, Feature Selection

Date: 27.09.2022

Plan for today

- Evaluation metrics
- Feature Selection

Evaluation metrics for binary classification

actual negative	True Negative	False Positive
actual positive	False Negative	True Positive
	predicted negative	predicted positive

$$\text{Accuracy} = \frac{\text{TP} + \text{TN}}{\text{TP} + \text{TN} + \text{FP} + \text{FN}}$$

Problems with Accuracy

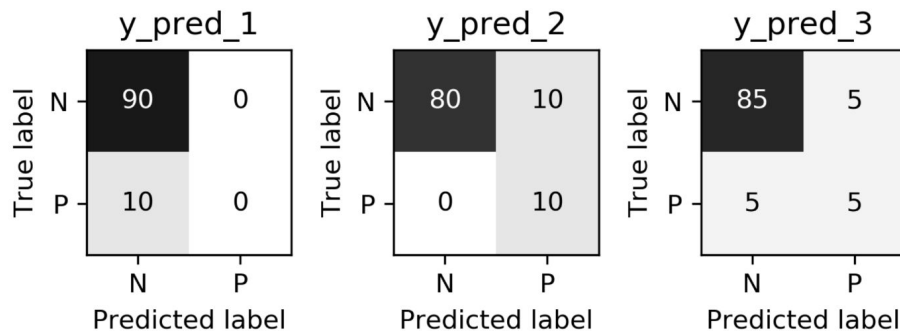
Data with 90% negatives:

```
from sklearn.metrics import accuracy_score
for y_pred in [y_pred_1, y_pred_2, y_pred_3]:
    print(accuracy_score(y_true, y_pred))
```

0.9

0.9

0.9



Evaluation Metrics for Classification

$$\text{Precision} = \frac{\text{TP}}{\text{TP} + \text{FP}}$$

Positive Predicted Value (PPV)

$$\text{Recall} = \frac{\text{TP}}{\text{TP} + \text{FN}}$$

Sensitivity, coverage, true positive rate.

$$F = 2 \frac{\text{precision} \cdot \text{recall}}{\text{precision} + \text{recall}}$$

Harmonic mean of precision and recall

Evaluation Metrics for Classification

```
classification_report(y_true, y_pred)
```

		y_pred_1			precision	recall	f1-score	support
True label	N	90	0	0	0.90	1.00	0.95	90
	P	10	0	1	0.00	0.00	0.00	10
		N		P				
		Predicted label						
						</		

Averaging Strategies

$$\text{macro } \frac{1}{|L|} \sum_{l \in L} R(y_l, \hat{y}_l)$$

$$\text{weighted } \frac{1}{n} \sum_{l \in L} n_l R(y_l, \hat{y}_l)$$

```
print("Weighted average: ", recall_score(y_test, y_pred_1, average="weighted"))  
print("Macro average: ", recall_score(y_test, y_pred_1, average="macro"))
```

Weighted average: 0.90

Macro average: 0.50

Balanced Accuracy

```
balanced_accuracy_score(y_t, y_p) == recall_score(y_t, y_p, average='macro')
```

$$\text{balanced_accuracy} = \frac{1}{2} \left(\frac{TP}{TP + FN} + \frac{TN}{TN + FP} \right)$$

- Always 0.5 for chance predictions
- Equal to accuracy for balanced datasets

Changing Thresholds

```
y_pred = rf.predict(X_test)
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
False	0.99	1.00	0.99	2732
True	0.90	0.56	0.69	64
accuracy			0.99	2796
macro avg	0.94	0.78	0.84	2796
weighted avg	0.99	0.99	0.99	2796

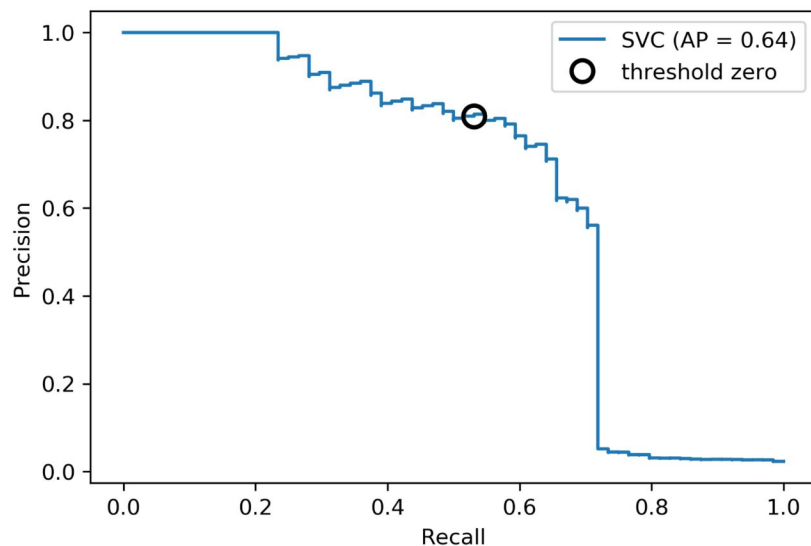
```
y_pred = rf.predict_proba(X_test)[: , 1] > .30
```

```
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
False	0.99	0.99	0.99	2732
True	0.71	0.64	0.67	64
accuracy			0.99	2796
macro avg	0.85	0.82	0.83	2796
weighted avg	0.99	0.99	0.99	2796

Precision-Recall Curve

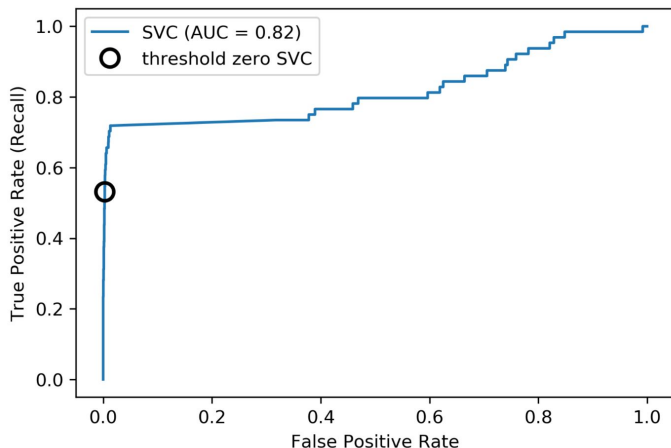
```
svc = make_pipeline(StandardScaler(), SVC(C=100, gamma=0.1))  
svc.fit(X_train, y_train)  
plot_precision_recall_curve(svc, X_test, y_test, name='SVC')
```



ROC Curve

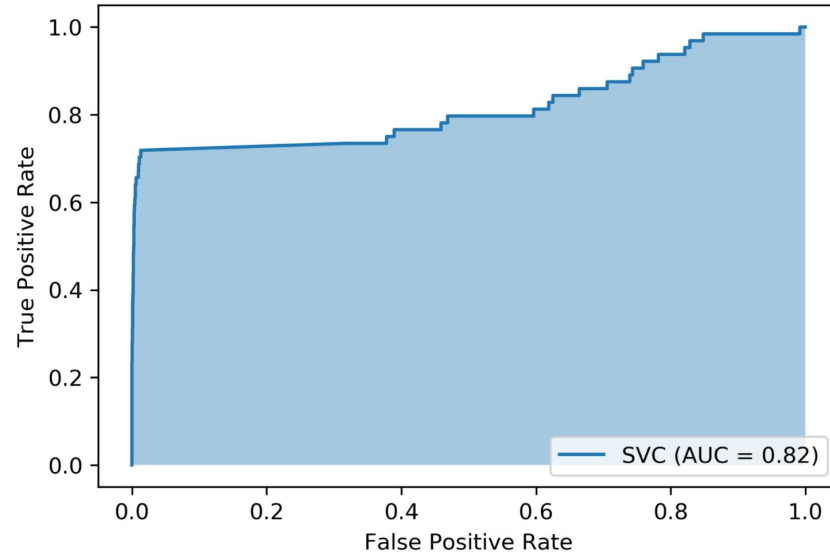
(Receiver Operating Characteristic)

```
plot_roc_curve(svc, X_test, y_test, name='SVC')
```



- True positive rate (recall)
- False Positive Rate (FPR)
 - Negative instances that are incorrectly classified as positive.
 - $1 - \text{True negative rate}$ (specificity)

Area Under ROC Curve (AUC)



- Always .5 for random/constant prediction

Summary of metrics for binary classification

- Threshold-based
 - (balanced) accuracy
 - precision , recall, f1
- Ranking
 - Average precision
 - ROC AUC

Picking metrics

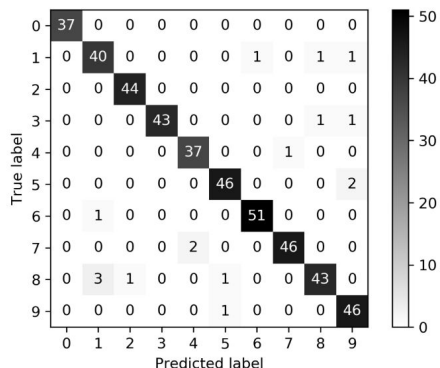
- Accuracy rarely what you want
- Problems are rarely balanced
- Find the right criterion for the task
- OR pick a substitute, but at least think about it
- Emphasis on recall or precision?
- Which classes are the important ones?

Multi-class classification

```
from sklearn.datasets import load_digits
from sklearn.metrics import accuracy_score

digits = load_digits()
# data is between 0 and 16
X_train, X_test, y_train, y_test = train_test_split(
    digits.data / 16., digits.target, random_state=0)
lr = LogisticRegression().fit(X_train, y_train)
pred = lr.predict(X_test)
print("Accuracy: {:.3f}".format(accuracy_score(y_test, pred)))
plot_confusion_matrix(lr, X_test, y_test, cmap='gray_r')
```

Accuracy: 0.964

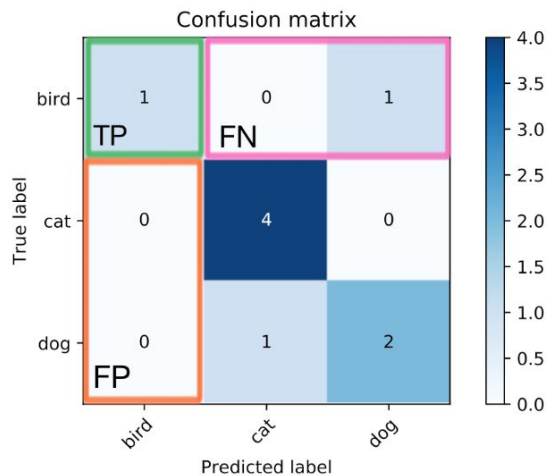


```
print(classification_report(y_test, pred))
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	37
1	0.91	0.93	0.92	43
2	0.98	1.00	0.99	44
3	1.00	0.96	0.98	45
4	0.95	0.97	0.96	38
5	0.96	0.96	0.96	48
6	0.98	0.98	0.98	52
7	0.98	0.96	0.97	48
8	0.96	0.90	0.92	48
9	0.92	0.98	0.95	47
accuracy			0.96	450
macro avg	0.96	0.96	0.96	450
weighted avg	0.96	0.96	0.96	450

Multiclass classification

Label	Predicted
cat	cat
cat	cat
cat	cat
cat	cat
dog	dog
dog	dog
dog	cat
bird	dog
bird	bird



	TP	FP	FN
bird	1	0	1
cat	4	1	0
dog	2	1	1
TOTAL	7	2	2

$$Precision_{birds} = \frac{TP_{birds}}{TP_{birds} + FP_{birds}} = \frac{1}{1 + 0} = 1$$

$$Precision_{cats} = \frac{TP_{cats}}{TP_{cats} + FP_{cats}} = \frac{4}{4 + 1} = 0.8$$

$$Precision_{dogs} = \frac{TP_{dogs}}{TP_{dogs} + FP_{dogs}} = \frac{2}{2 + 1} = 0.667$$

Multiclass classification

- **Micro-averaged:** all samples equally contribute to the final averaged metric
- **Macro-averaged:** all classes equally contribute to the final averaged metric
- **Weighted-averaged:** each classes contribution to the average is weighted by its size

	TP	FP	FN	Precision	Number of samples
bird	1	0	1	1	2
cat	4	1	0	0.8	4
dog	2	1	1	0.667	3
TOTAL	7	2	2		

$$\text{Micro-averaged Precision} = \frac{TP_{total}}{TP_{total} + FP_{total}} = \frac{7}{7 + 2} = 0.7777$$

$$\text{Macro-averaged Precision} = \frac{1}{3} Precision_{birds} + Precision_{cats} + Precision_{dogs} = \frac{1}{3} (1 + 0.8 + 0.6666) = 0.8222$$

$$\text{Weighted-averaged Precision} = \frac{Precision_{birds} * N_{birds} + Precision_{cats} * N_{birds} + Precision_{dogs} * N_{birds}}{\text{Total number of samples}} = \frac{1 * 2 + 0.8 * 4 + 0.6666 * 3}{2 + 4 + 3} = 0.8$$

Feature Selection

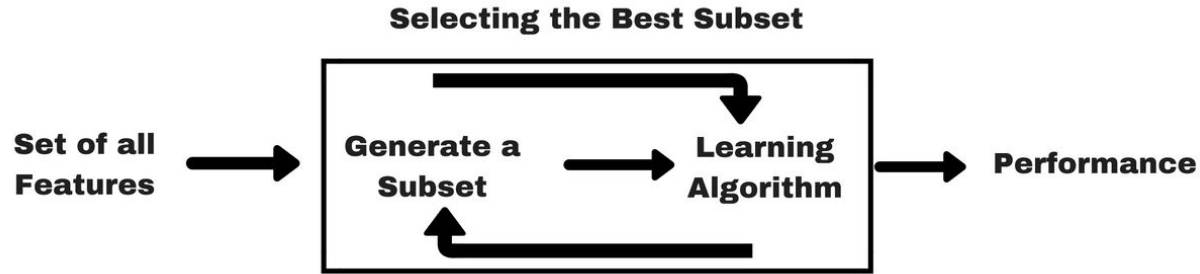
- Filtering-based feature selection
- Wrapper-based feature selection
- Embedded feature selection
 - Regularization
- Dimensionality Reduction
 - Do not select features, instead construct new features that are effectively represent combinations original features
- Motivation
 - Avoid overfitting
 - Faster prediction and training
 - Less storage for model
 - More interpretable model

Filter Based Feature Selection



- Variance-based: 0 variance or mostly constant
- Covariance-based: remove correlated features (or pearson corr.)
- Chi-Square test: a statistical test that compares the frequencies of a term between different classes

Wrapper Based Feature Selection



- Forward Selection
- Backward Elimination
- Recursive Feature Elimination

Forward Feature Selection

- **Forward Feature Selection:** is an iterative method in which we start having no feature in the model. In each iteration, we keep adding the feature which best improves our model till an addition of a new variable does not improve the performance of the model.

Given: feature set $\{X_1, \dots, X_n\}$, training set D , learning method L

$F \leftarrow \{ \}$

while score of F is improving

for $i \leftarrow 1$ to n do

if $X_i \notin F$


$G_i \leftarrow F \cup \{X_i\}$

$Score_i = \text{Evaluate}(G_i, L, D)$

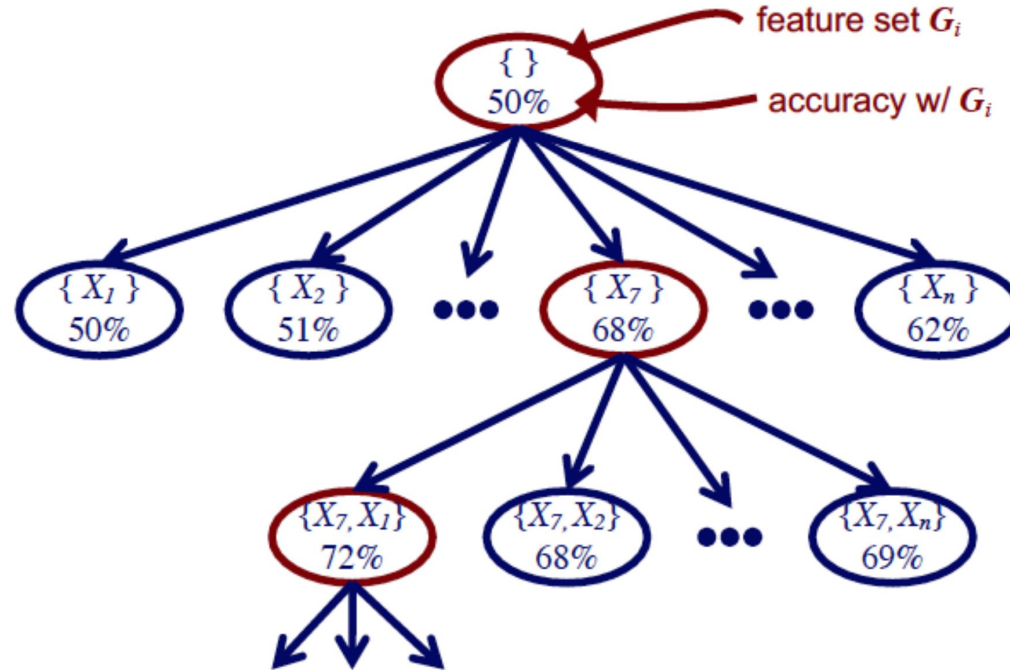
$F \leftarrow G_b$ with best $Score_b$

return feature set F

scores feature set G by learning model(s) with L and assessing its (their) accuracy



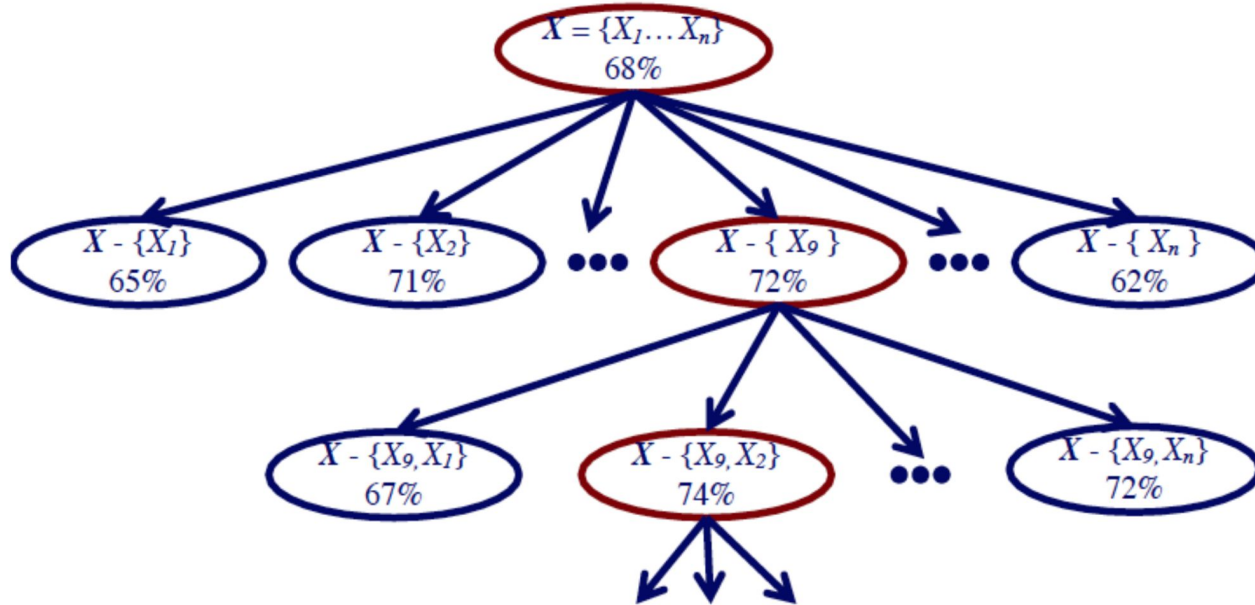
Forward Feature Selection



Backward Feature Elimination

- **Backward Feature Elimination:** In backward elimination, we start with all the features and removes the least significant feature at each iteration which improves the performance of the model. We repeat this until no improvement is observed on removal of features.

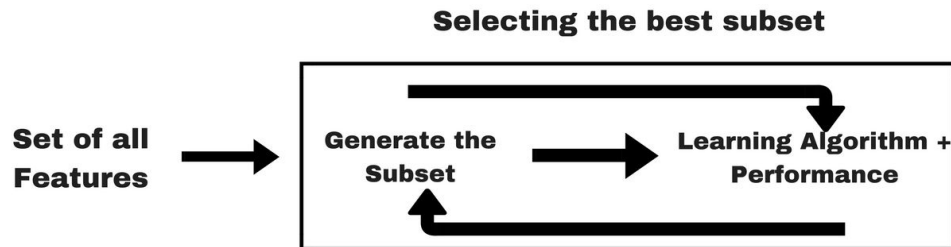
Backward Feature Elimination



Forward Selection vs Backward Elimination

- Forward Selection
 - Efficient for choosing a small subset of the features
 - Misses features whose usefulness requires other features (feature synergy)
- Backward elimination
 - Efficient for discarding a small subset of the features
 - Preserves features whose usefulness requires other features

Embedding Based Feature Selection



- **Regularization**
 - Lasso Regression
 - Ridge Regression
- **Model-Based Selection**
 - Feature importance ()

Difference between Filter and Wrapper Methods

- Filter methods measure the relevance of features by their correlation with dependent variable, while wrapper methods measure the usefulness of a subset of feature by actually training a model on it.
- Filter methods are much faster compared to wrapper methods.
- Filter methods use statistical methods for evaluation of a subset of features while wrapper methods use cross validation.
- Filter methods might fail to find the best subset of features in many occasions.
- Using the subset of features from the wrapper methods make the model more prone to overfitting as compared to using subset of features from the filter methods.

Next Class:

Logistic Regression, Neural Network